

# **A short introduction to ER & SQL**



Bipin C. DESAI

Pl. see: <https://users.encs.concordia.ca/~bcdesai/CopyForward.pdf>



# Database Languages

- ❖ A Database Management System (DBMS) provides two types of languages; they may also be viewed as components of the DBMS language:
  - ◆ Data Definition Language (DDL)
    - Language (notation) for *defining* and modifying a database schema
    - It includes syntax for declaring tables, indexes, views, constraints, etc.)
  - ◆ Data Manipulation Language (DML)
    - Language for *accessing* and *manipulating* the data (organized/stored according to the appropriate data model)

# Query Languages

- ❖ Theoretical:
  - Relational Algebra, Relational Calculus, Datalog
- ❖ Commercial: SQL
- ❖ First there were two: SEQUEL (Ingres) and SQL(R)
- ❖ SQL developed originally at IBM in 1976
  - ◆ First standard: SQL-86
  - ◆ Second standard: SQL-92
  - ◆ Latest standard: SQL-99, or SQL3,
    - SQL3 standard has over 1,000 pages of document
- ❖ SQL is the de-facto standard for RDBMS
- ❖ The SQL query language components:
  - ◆ DDL (e.g., create)
  - ◆ DML (e.g., select, insert, update, delete)

# Simple SQL Queries

A SQL query has a form:

**SELECT** . . .

**FROM** . . .

**WHERE** . . . ;

The **SELECT** clause defines the schema of the result

The **FROM** clause gives the source relation(s) of the query

The **WHERE** clause is one or more predicates to ‘select’ the tuples of interest.

The query result is a relation and it is **unnamed**.

# Example “theoretical” SQL Query

Relation schema:

**Course** (Cno, Cname, credits)

Query in natural language (English):

Find all the courses stored in the database

Query in SQL:

**SELECT \***

**FROM** Course;

Here the “\*” in “**SELECT \***” means **all** attributes in the **relation(s)** involved.

# More Examples SQL Query

❖ Relation schema:

**Movie** (title, year, length, filmType)

Query in natural language (English):

Find the titles of all movies stored in the database

Query in SQL:

**SELECT** title

**FROM** Movie;

Relation schema:

**Student** (SID, FirstName, LastName, Address, GPA)

Query in natural language (English):

Find the SID of every student whose GPA is greater than 3

Query in SQL:

**SELECT** SID

**FROM** Student

**WHERE** GPA > 3;

# The “WHERE” clause

The expressions that may follow **WHERE** are conditions

Standard comparison operators  $\theta$  includes  $\{ =, <>, <, >, <=, >= \}$

The values that may be compared include constants and attributes of the relation(s) mentioned in **FROM** clause

Simple expression

**A**  $\theta$  Value

**A**  $\theta$  **B**

Where **A**, **B** are attributes and

$\theta$  is a comparison operator

We may also apply the usual arithmetic operators, +, -, \*, /, etc. to numeric values before comparing them

**(year - 1930) \* (year - 1930) < 100**

The result of a comparison is a Boolean value **TRUE** or **FALSE**

Boolean expressions can be combined by the logical operators **AND**, **OR**, and **NOT**

- ❖ Relation schema: **Movie** (title, year, length, filmType)
- ❖ Query: Find the titles of all color movies produced in 1950
- ❖ Query in SQL:  
**SELECT** title  
**FROM** Movie  
**WHERE** filmType = 'color' **AND** year = 1950;
- ❖ Query: Find the titles of color movies that are either made after 1970 or are less than 90 minutes long
- ❖ Query in SQL:  
**SELECT** title  
**FROM** Movie  
**WHERE** (year > 1970 **OR** length < 90) **AND** filmType = 'color';

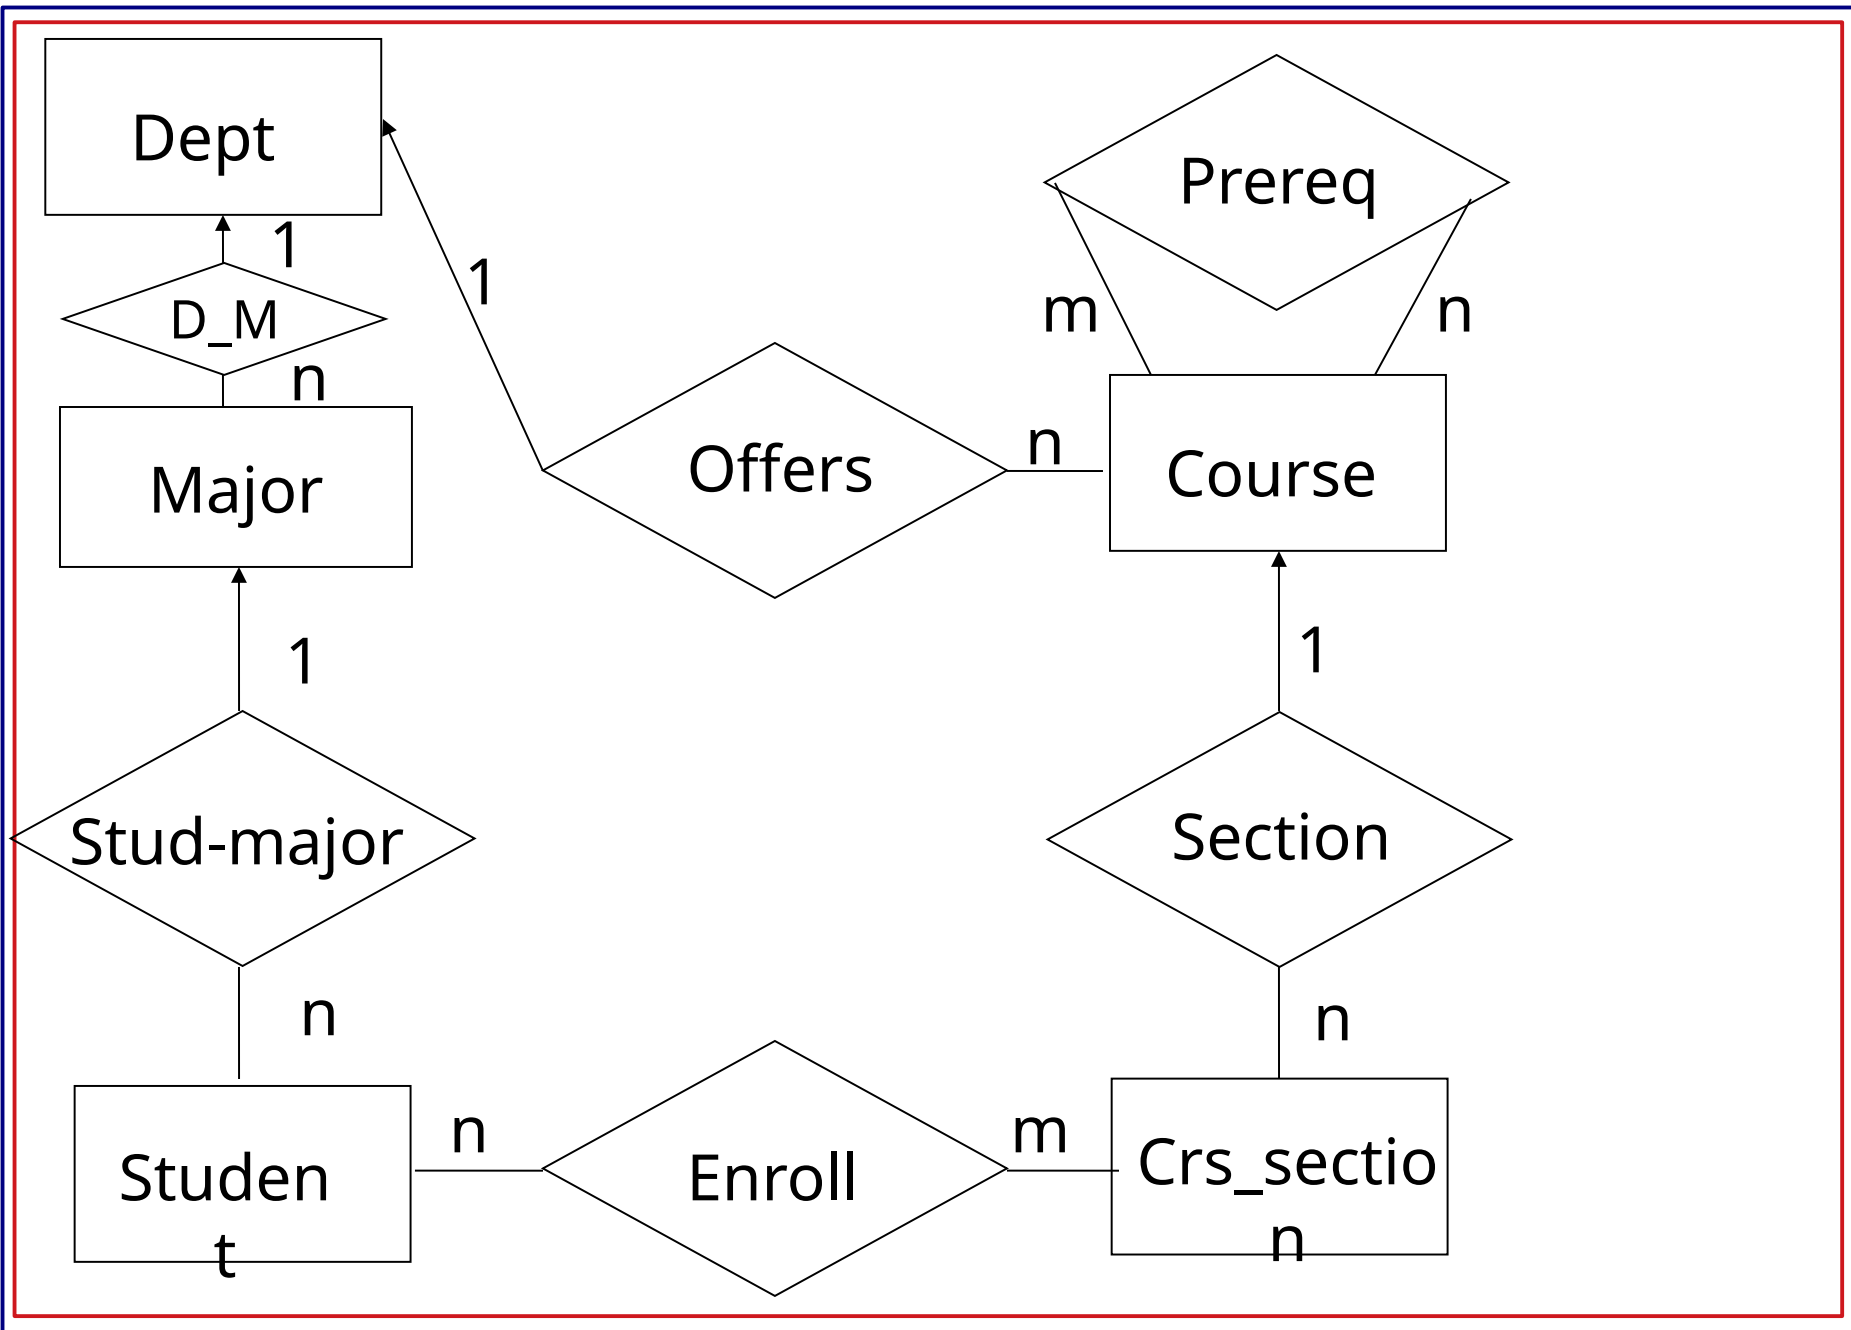
Note the precedence rules, when parentheses are absent:

**AND** takes precedence over **OR**,  
and **NOT** takes precedence over both




# **An example of using SQL from E-R to RDBMS**

**© Bipin C. Desai**



SQL DDL example:



create table student

(SID NUMBER(3) primary key not null,

SNAME VARCHAR2(20),

MAJOR CHAR(4),

YEAR NUMBER(1),

BDATE DATE)

/

insert into student values

(8,'Brenda','COMP','2','13-AUG-77');

(8,'Brenda','COMP','1','1977-08-13');

insert into student values

(10,'Mary','ENGL','1','13-MAY-80');

(10,'Mary','ENGL','1','1980-5-13');

insert into student values

(13,'Keily','SENG','4','12-AUG-80');

insert into student values


(14,'Jack','CSAP','1','12-FEB-77');

# Many to one relationship

A department offers many courses

A given course can be offered by only one department

There is no standard regarding the direction of arrow for the “one” entity.



*Alternate ways of representing  
The “one” of the many-to-one  
relationship; arrow either pointing to entity on  
the “one side” or pointing to the relationship*

```
create table course  
(COURSE_NAME CHAR(20),  
  COURSE_NUMBER CHAR(8) primary key NOT  
  NULL,  
  CREDIT_HOURS NUMBER(2),  
  OFFERING_DEPT CHAR(4))  
  tablespace TUTOR pctfree 5
```

/

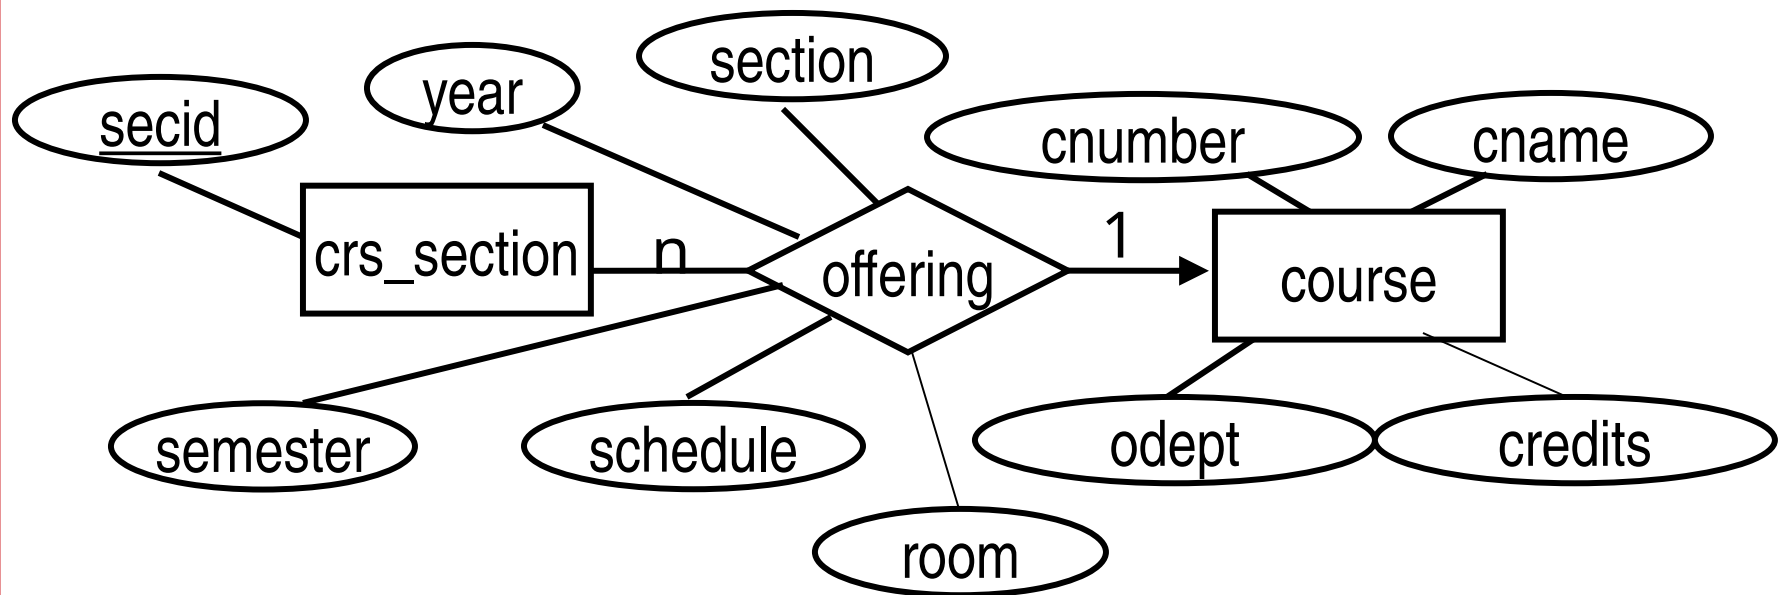
Note how the relationship  
w/o an attribute is  
“merged” with one of the  
entity

**insert into course values  
('C++','COMP248',3,'COMP');**

**insert into course values  
('DATA STRUCTURES ','COMP352',3,'COMP');**

**insert into course values  
('OPERATING SYSTEMS','COMP346',4,'COMP');**

**insert into course values  
('DATABASE','COMP353',4,'COMP');**



**NB: Here course section (crssection) is really a “weak” entity; However, in most cases it is promoted a “strong” entity by introducing an identifying key attribute section ID (secid)**



```
create table crs_section
```

```
(SECID NUMBER(6) primary key NOT NULL,
```

```
COURSE_NUM CHAR(8),
```

```
SECTION CHAR(2),
```

```
SEMESTER CHAR(4),
```

```
YEAR CHAR(4),
```

```
SCHEDULE CHAR(10),
```

```
ROOM CHAR(7))
```

```
tablespace TUTOR pctfree 2
```

```
/
```

Note: We have replaced an  
entity and the relationship  
With a single relation

insert into crs\_section values

(85,'COMP352','A','FALL', '1998','TH16001715','H123');

insert into crs\_section values

(90,'COMP353','B','FALL','1999','MW08451000','H631');

```
create table enrolment
(STUDENT_NUMBER NUMBER(3) not null,
SECTION_ID NUMBER(6) not null,
GRADE CHAR(1),
primary key(student_number, section_id))
tablespace TUTOR pctfree

/
insert into enrolment values(8,85,null);
insert into enrolment values(10,90,null);
insert into enrolment values(8,90,null);
insert into enrolment values(14,90,null);
```

**Find details of studs. taking a course offered by the “DISC” dept.**

```
select s.SID, s.SNAME, s.MAJOR, s.YEAR, s.BDATE  
from student s, dept d, course c, crs_section r, enrolment e  
where c.ODEPT=d.CODE and  
      r.COURSE_NUM=c.CNUMBER and  
      r.SECID=e.SECTION_ID and  
      e.STUDENT_NUMBER = s.SID and  
      d.CODE= 'DISC';
```

<b>SID</b>	<b>SNAME</b>	<b>MAJOR</b>	<b>YEAR</b>	<b>BDATE</b>
-----				
<b>1 , 4</b>	<b>Jack</b>	<b>CSAP</b>	<b>1</b>	<b>12-FEB-77</b>

# **More examples of using E-R modeling**

**© Bipin C. Desai**

Professors have a SIN, a name, an age, a rank, and a research specialty.

Projects have a project number, a sponsor name (e.g., NSERC),  
a starting date, an ending date, and a budget.

Graduate students have a SIN, a name, an age, and a degree program  
(e.g., M.S. or Ph. D.).

Each project is managed by a professor (principal investigator).

Each project is worked on by one or more professors (co-investigators).

Professors can manage and/or work on multiple projects.

Each project is worked on by one or more graduate students (research assistants).

When a graduate student works on a project, is supervised by a  
participating professor.

Graduate students can work on multiple projects.

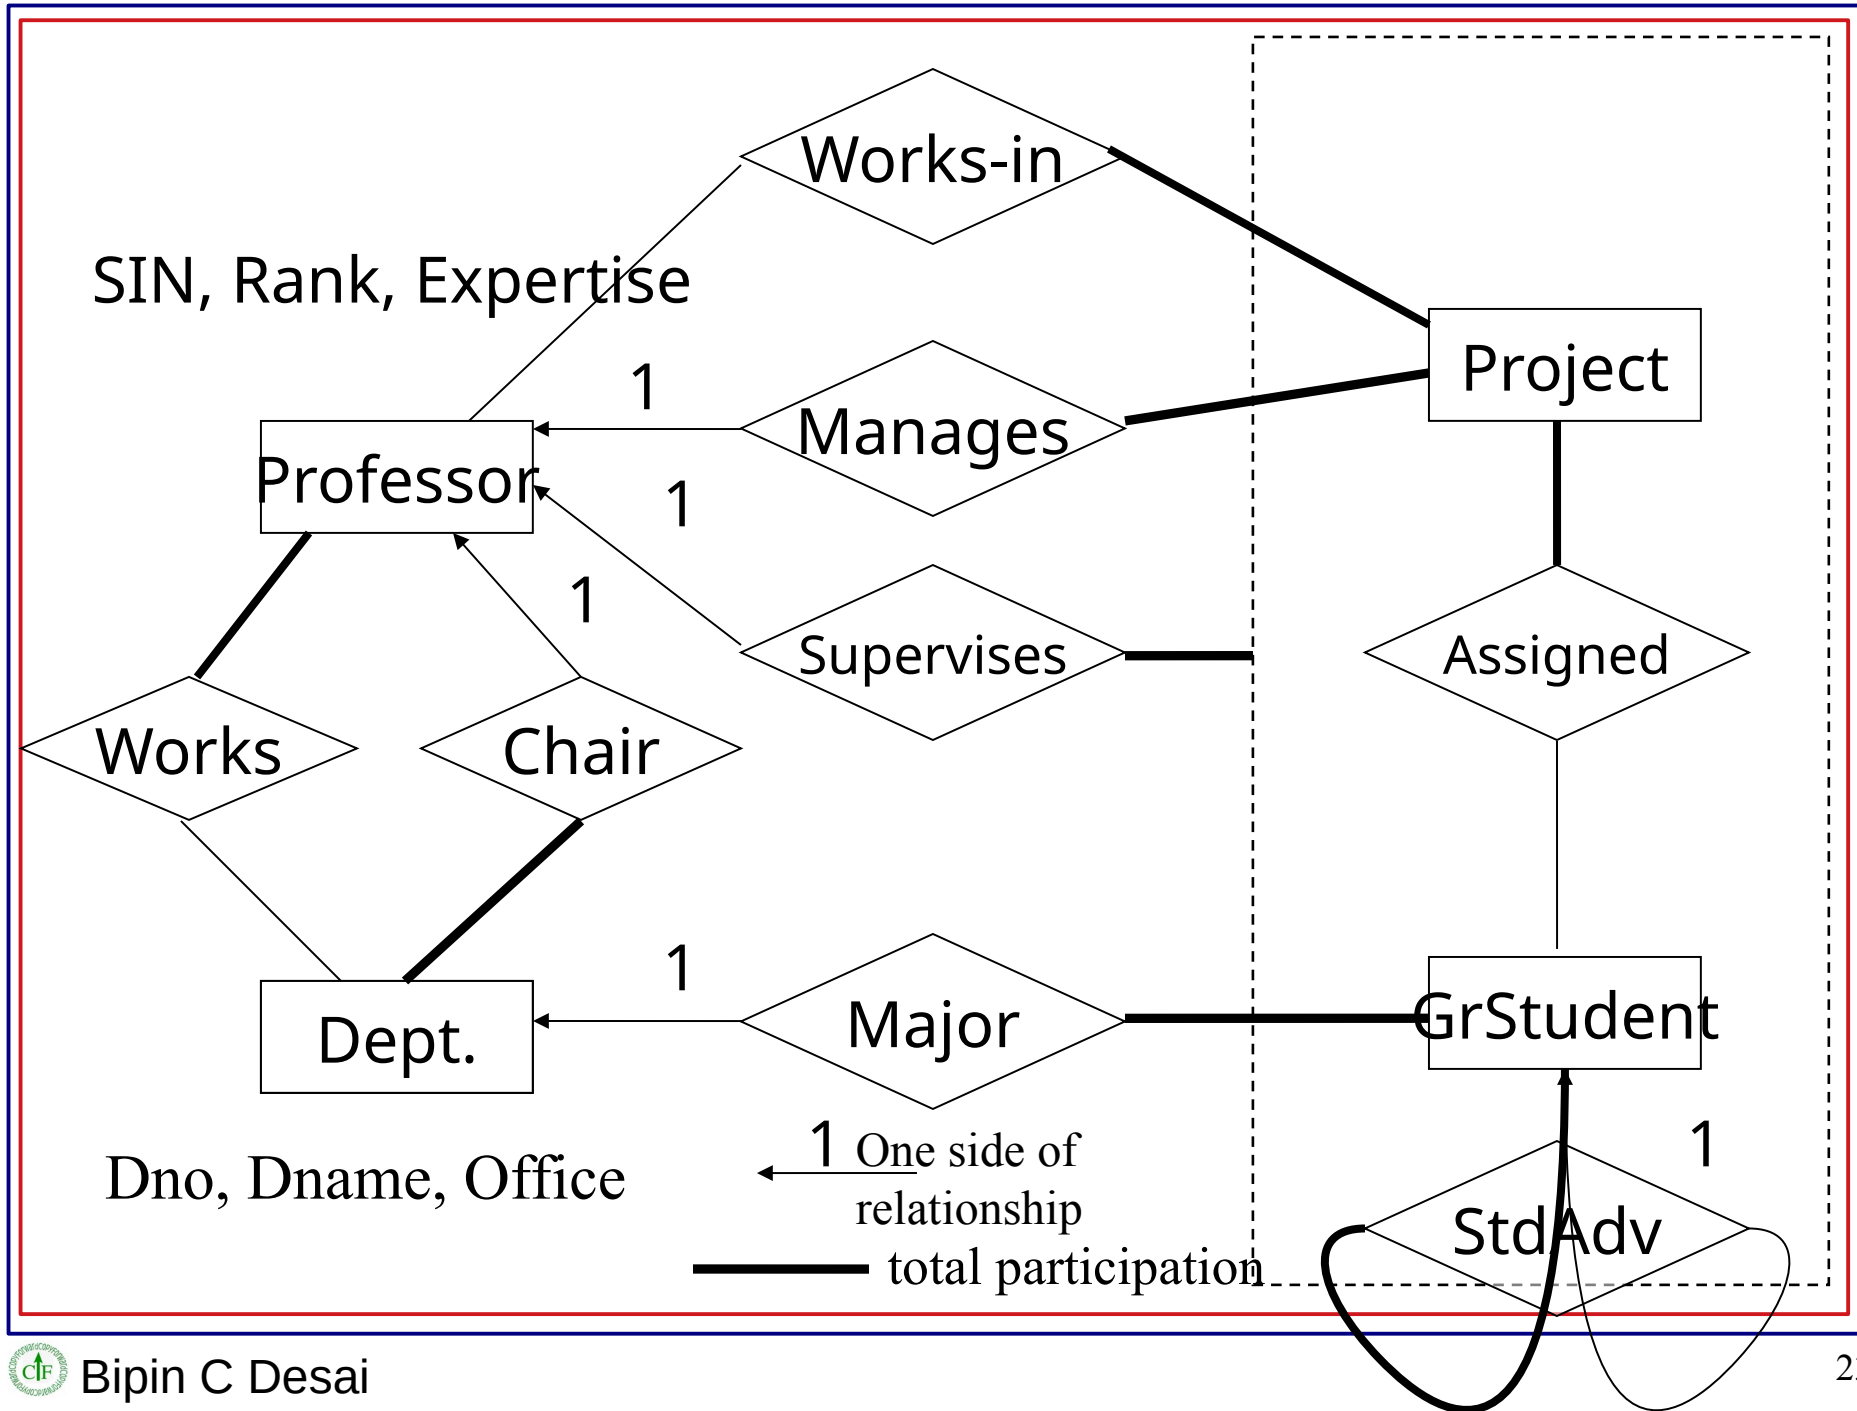
Departments have a department number, a department name, and a main office.

Departments have a chairman who runs the department.

Professors work in one or more departments(%time)


Graduate students have one major department for their degree.

Each graduate student has another, more senior graduate student  
(student advisor)

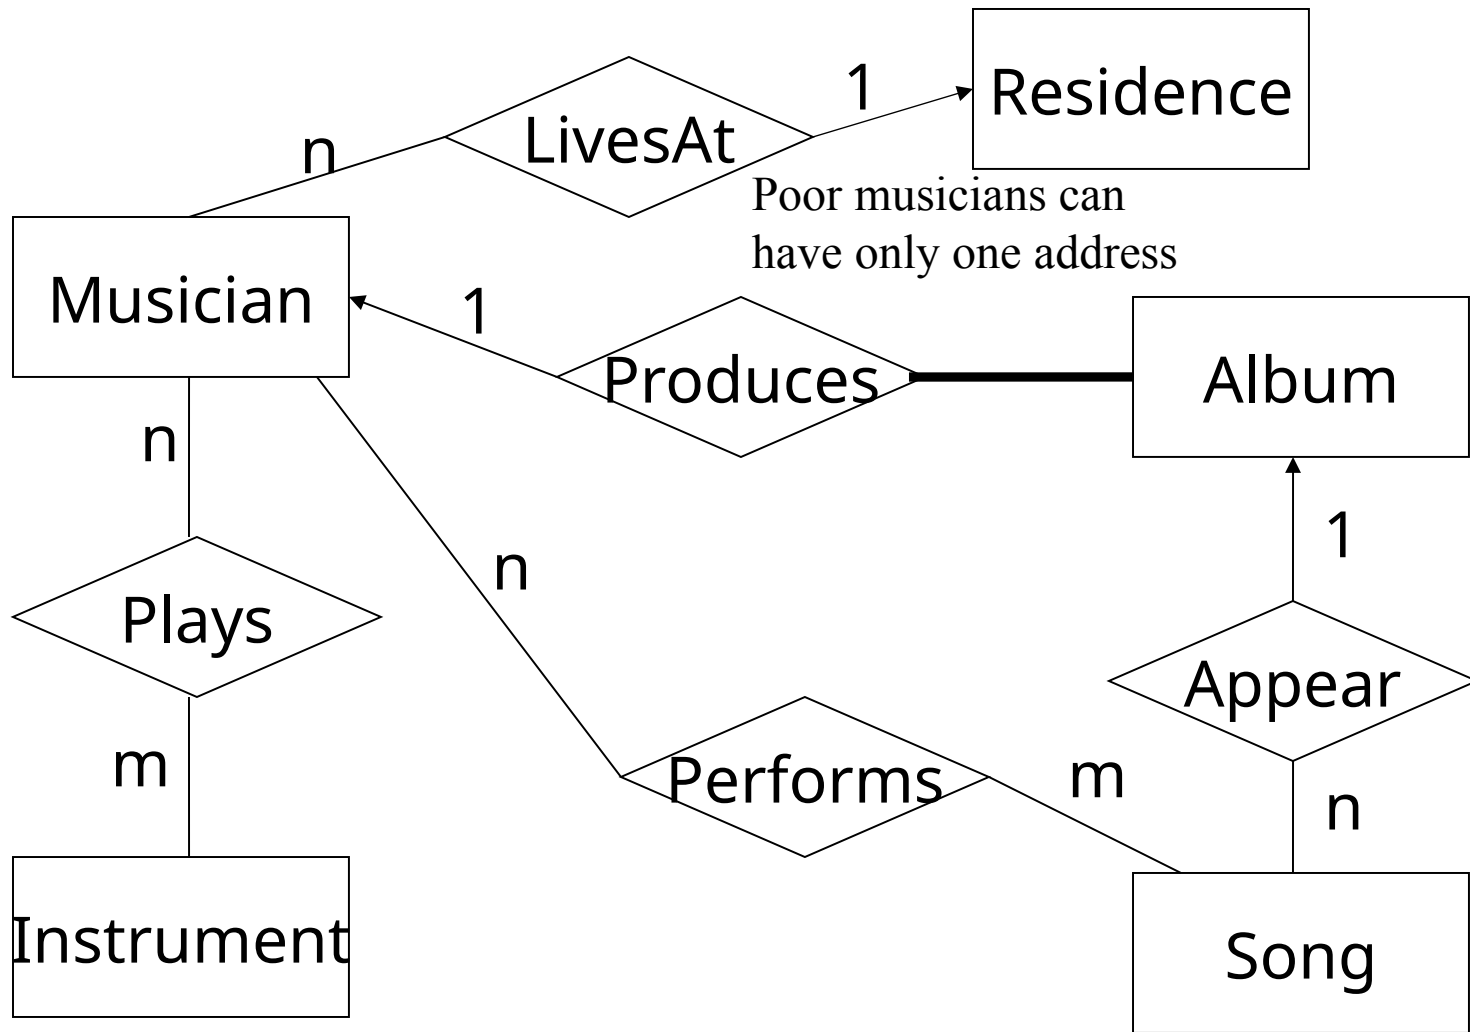


In a company database, you need to store information about employees (SIN, salary and phone), departments (dno, dname, budget), and children of employees (with name and age as attributes). Employees work in departments; Each employee works in only one department; A department could have many employees; Each department is managed by an employee; Each department has only one manager(an employee); A manager could manage many departments A child can only be identified by name An employee has only one child with a given name only one parent can declare a child as a dependent










Each musician that records at Notown has an SIN, a name,  
an address, and a phone number.  
Musicians often share the same address,  
no address has more than one phone.  
Each instrument that is used in songs recorded at Notown has a name and  
a musical key  
Each album that is recorded on has a title, a copyright date, and  
an album identifier.  
Each song recorded has a title and an author.  
Each musician may play several instruments, and  
a given instrument may be played by several musicians.  
Each album has a number of songs on it,  
but no song may appear on more than one album.  
Each song is performed by one or more musicians, and  
a musician may perform a number of songs.  
For each album, there is exactly one musician that acts as its producer.  
A musician may produce several albums.



Notes: Since a songs must appear on only one album, Appear is a many to one relationship. Similarly for Produces. Album requires total participation in Produces. Some songs may not be recorded and there may be some instruments that nobody can play!




Total participation of entry in relationship




Alternate E-R notations

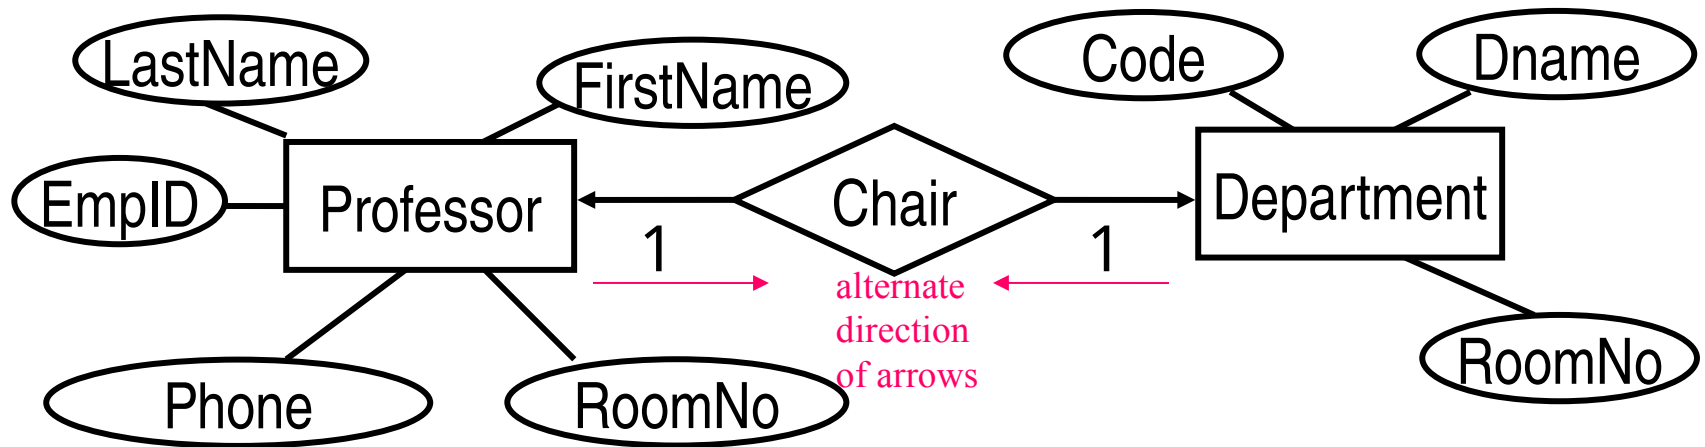
ISA: specialization or generalization



Total generalization



- ❖ A **one-one** relationship between **Department** and **its chair** ( a dept. has one chair and a prof. is the chair of at most one dept) is represented by
- ◆ arrows pointing to both Department and Professor or
  - ◆ indicated by a line with the number 1 on it.
  - ◆ Sometimes the arrow is in the opposite direction(pointing to the diamond)




**ODL** allows only **binary** relationships, i.e., relationships involving two classes.

**E/R** model makes it convenient to define **(n-ary)** relationships – relationships involving n entity sets

A **n-ary** relationship in an **E/R** diagram is represented by lines from the relationship (diamond) to each of the participating entity sets (rectangles).

# N-ary Relationships



Multiplicity of this ternary relationship: n to n to n

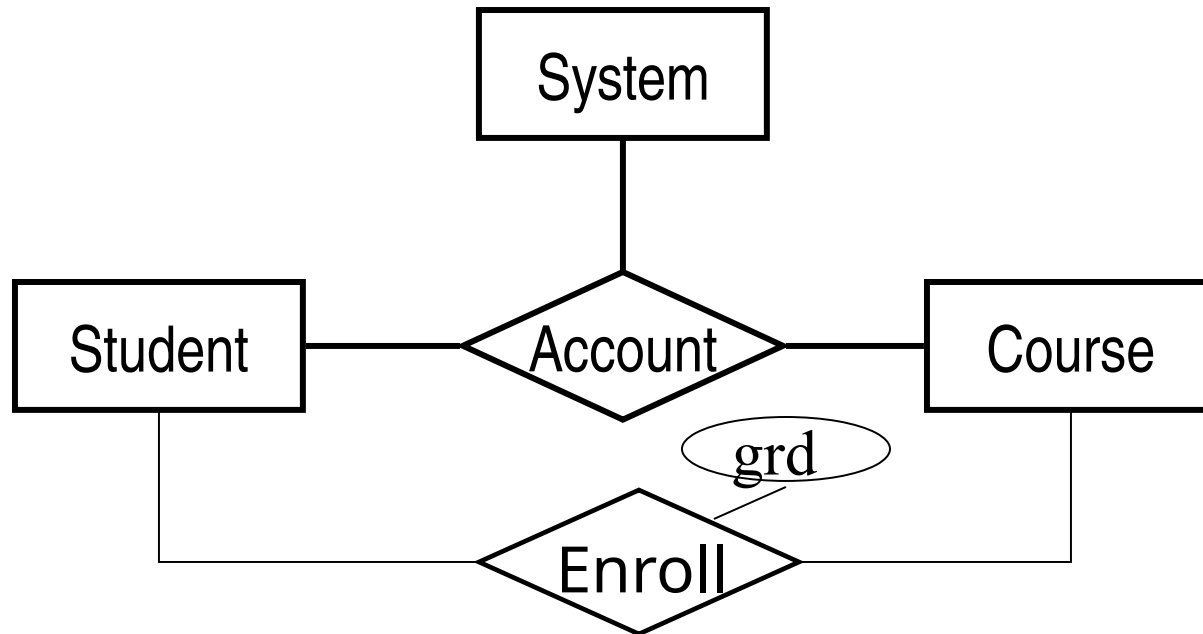
Enroll(sid(fk → Student), cno(fk → Course), sys(fk → System), grd)

What is the problem here?

What is the schema for Enroll?

How is the n:n:n relationship mapped to a relation(table)?

# N-ary Relationships




Multiplicity of this ternary relationship: n to n to n

Account( sid(fk → Student), cno(fk → Course), sys(fk → System), uid)

Enroll(sid(fk → Student), cno(fk → Course), grd)




# N-ary Relationships




Multiplicity of this N-ary relationship:  $n_1$ ----  $n_i$  ----  $n_n$

Any of this  $n_i$  could be 1


Any of this could be multiple



If a binary relationship between two entity sets is 1-to-1,  
- the primary key of the relationship is the key of the entity from either of the '1' side, the other side is a foreign key (*would be unique*)




If a binary relationship between two entity sets is 1-to-many,  
- the primary key of the relationship is from the 'm' side, the '1' side is the foreign key (*would be unique*)



If a binary relationship between two entity sets is m-to-n,  
- the primary key is composite, consisting of the primary key of the entities from each side of the relationship

# N-ary Relationships



Multiplicity of this ternary relationship: 1 to 1 to n

Offspring(fid (fatherID  $\rightarrow$  Person(ID)), motherID (fk  $\rightarrow$  Person(ID)),  
ChildID (fk  $\rightarrow$  Person(ID)))

**Who here is the father and mother?**

**What is the key?**

```

create table person(
ID number primary key,
gender char(1),
DOB date);
insert into person values(1, 'M', '11-Jan-1900');
insert into person values(2,'F', '11-Jan-1902');
insert into person values(121,'M', '11-Jan-1925');
insert into person values(122,'F', '11-Jan-1927');
insert into person values(3,'M', '11-Jan-1901');
insert into person values(4,'F', '11-Jan-1903');
insert into person values(341,'M', '11-Jan-1926');
insert into person values(342,'F', '11-Jan-1928');
insert into person
        values(1213421,'M', '11-Jan-1948');
insert into person
        values(1213422,'F', '11-Jan-1950');

```


**Could two tuples exist in offspring with the same cid???**

```

create table offspring(
fid number, mid number,
cid number primary key,
foreign key (fid)
        references person(id),
foreign key (mid)
        references person(id),
foreign key (cid)
        references person(id));
insert into offspring values(1,2,121);
insert into offspring
        values(1,2,122);
insert into offspring
        values(3,4,341);
insert into offspring
        values(3,4,342);
insert into offspring
        values(121,342, 1213421);
insert into offspring
        values(121, 342, 1213422);

```

# Replacing a ternary relation by a binary relation



**What is the schema for Offspring here?**

**What is a possible inconsistency problem?**

**Who is the father, mother??**

Offspring (IDC, IDF, IDM, DOB)

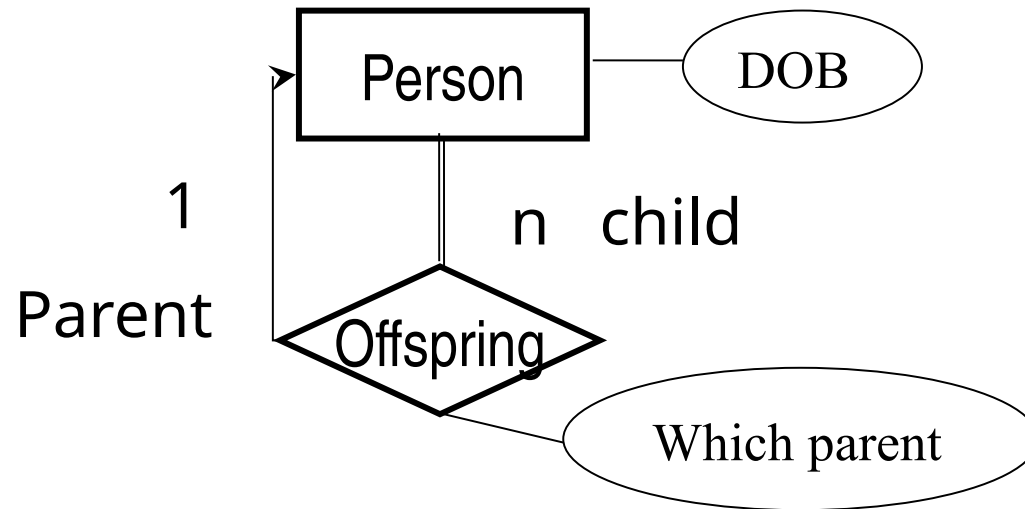
Father(IDC, IDF, DOB)

Mother(ID1, IDM, DOB)

- *not the same ER*
- *duplication of DOB*
- *Composite key*

# Replacing a ternary relation by a binary relation

## --- an alternate *non-normal form*



**What is the schema for Offspring?**  
**Is there a duplication problem?**  
**What is the primary key?**

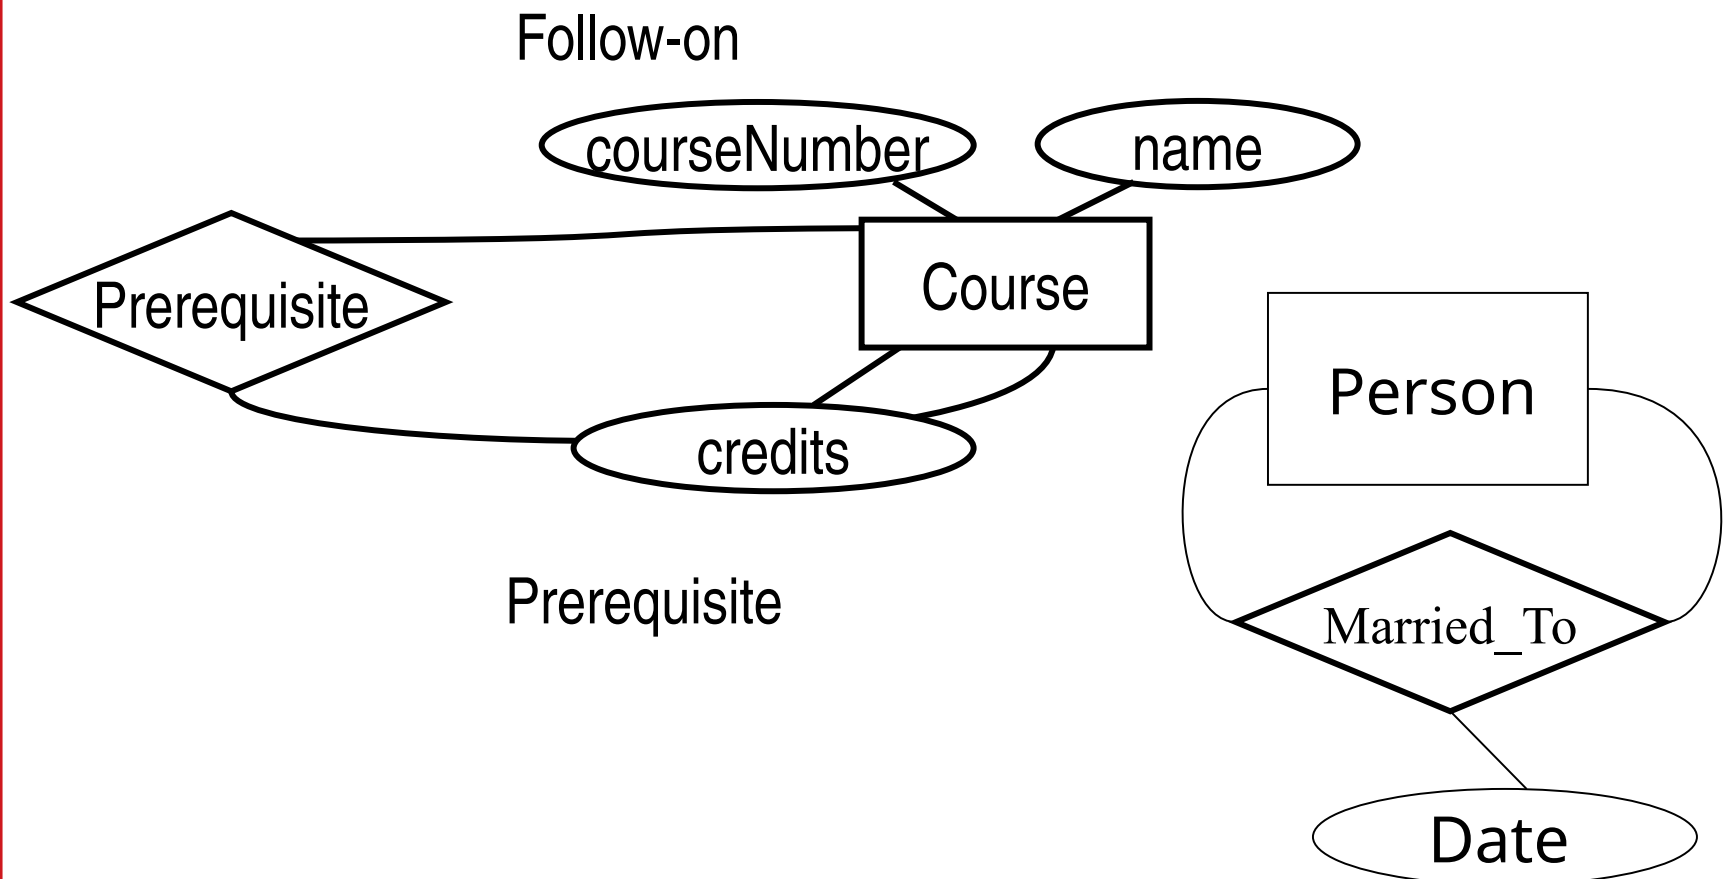
Multivalued attribute

121	1	Father
	2	Mother
122	1	Father
	2	Mother

# Roles in Relationships

- It is possible that one entity set appears **two** or **more** times in a relationship
- Suppose, we want to capture relationship between two **courses**, one of which is the **pre-requisite/follow-on** of the other

*Each line* to the entity set represents a different **role** that the entity set plays in the relationship

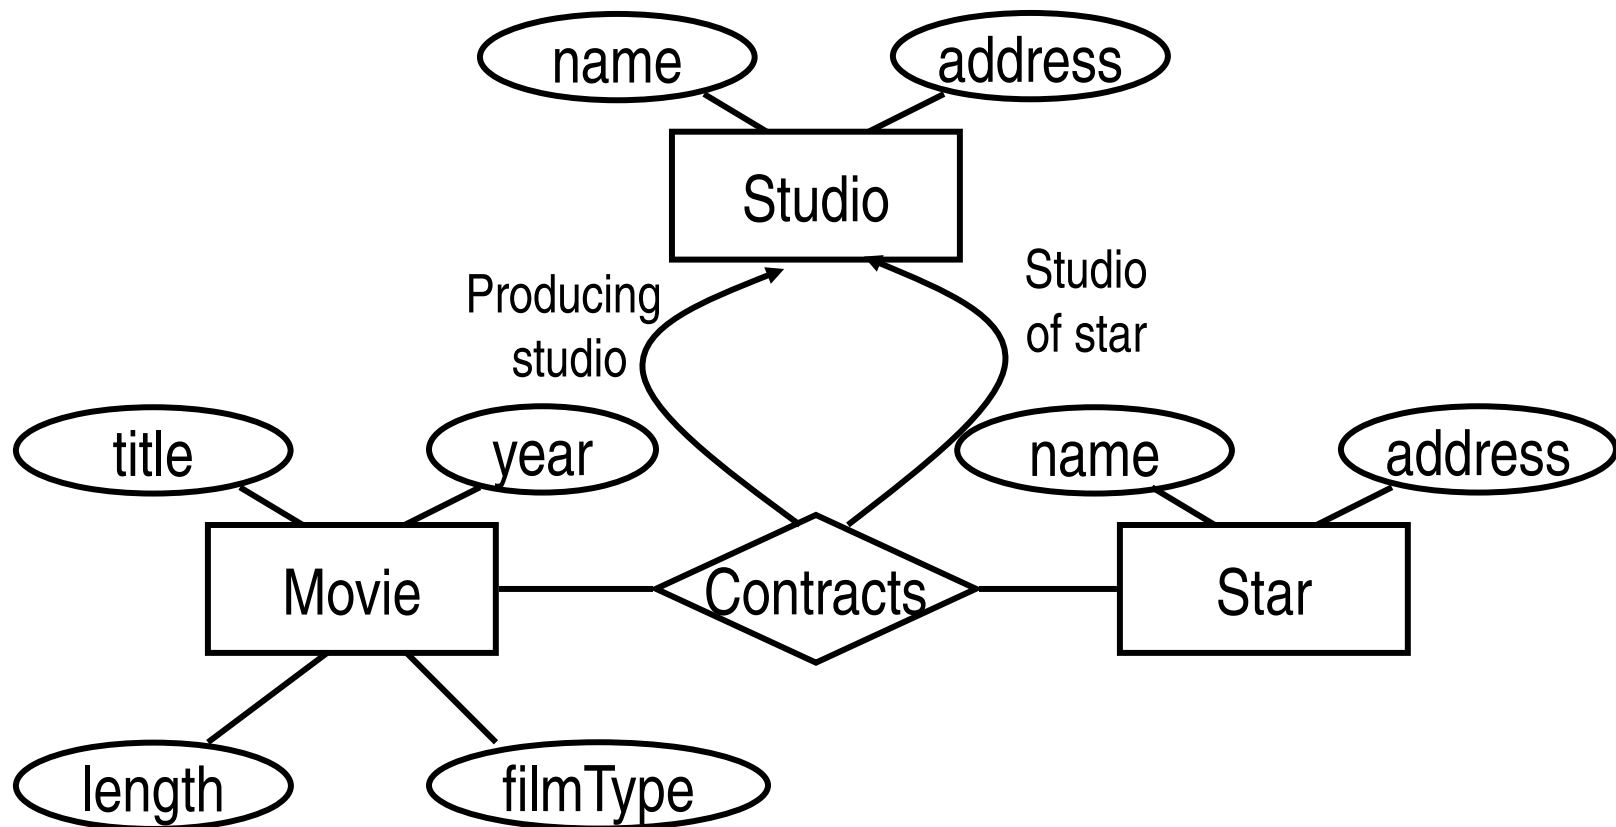




```
create table prereq  
(COURSE_Number CHAR(8),  
PREREQ CHAR(8),  
primary key (course_number, prereq))  
tablespace TUTOR pctfree 2  
/
```

```
insert into prereq values('COMP353','COMP352');  
insert into prereq values('COMP353','COMP346');  
insert into prereq values('COMP352','COMP248');
```

Suppose, each star is under contract with a single studio  
The studio of the star may enter into a contract with another studio to allow that star to act in a particular movie



# Converting n-ary relationship

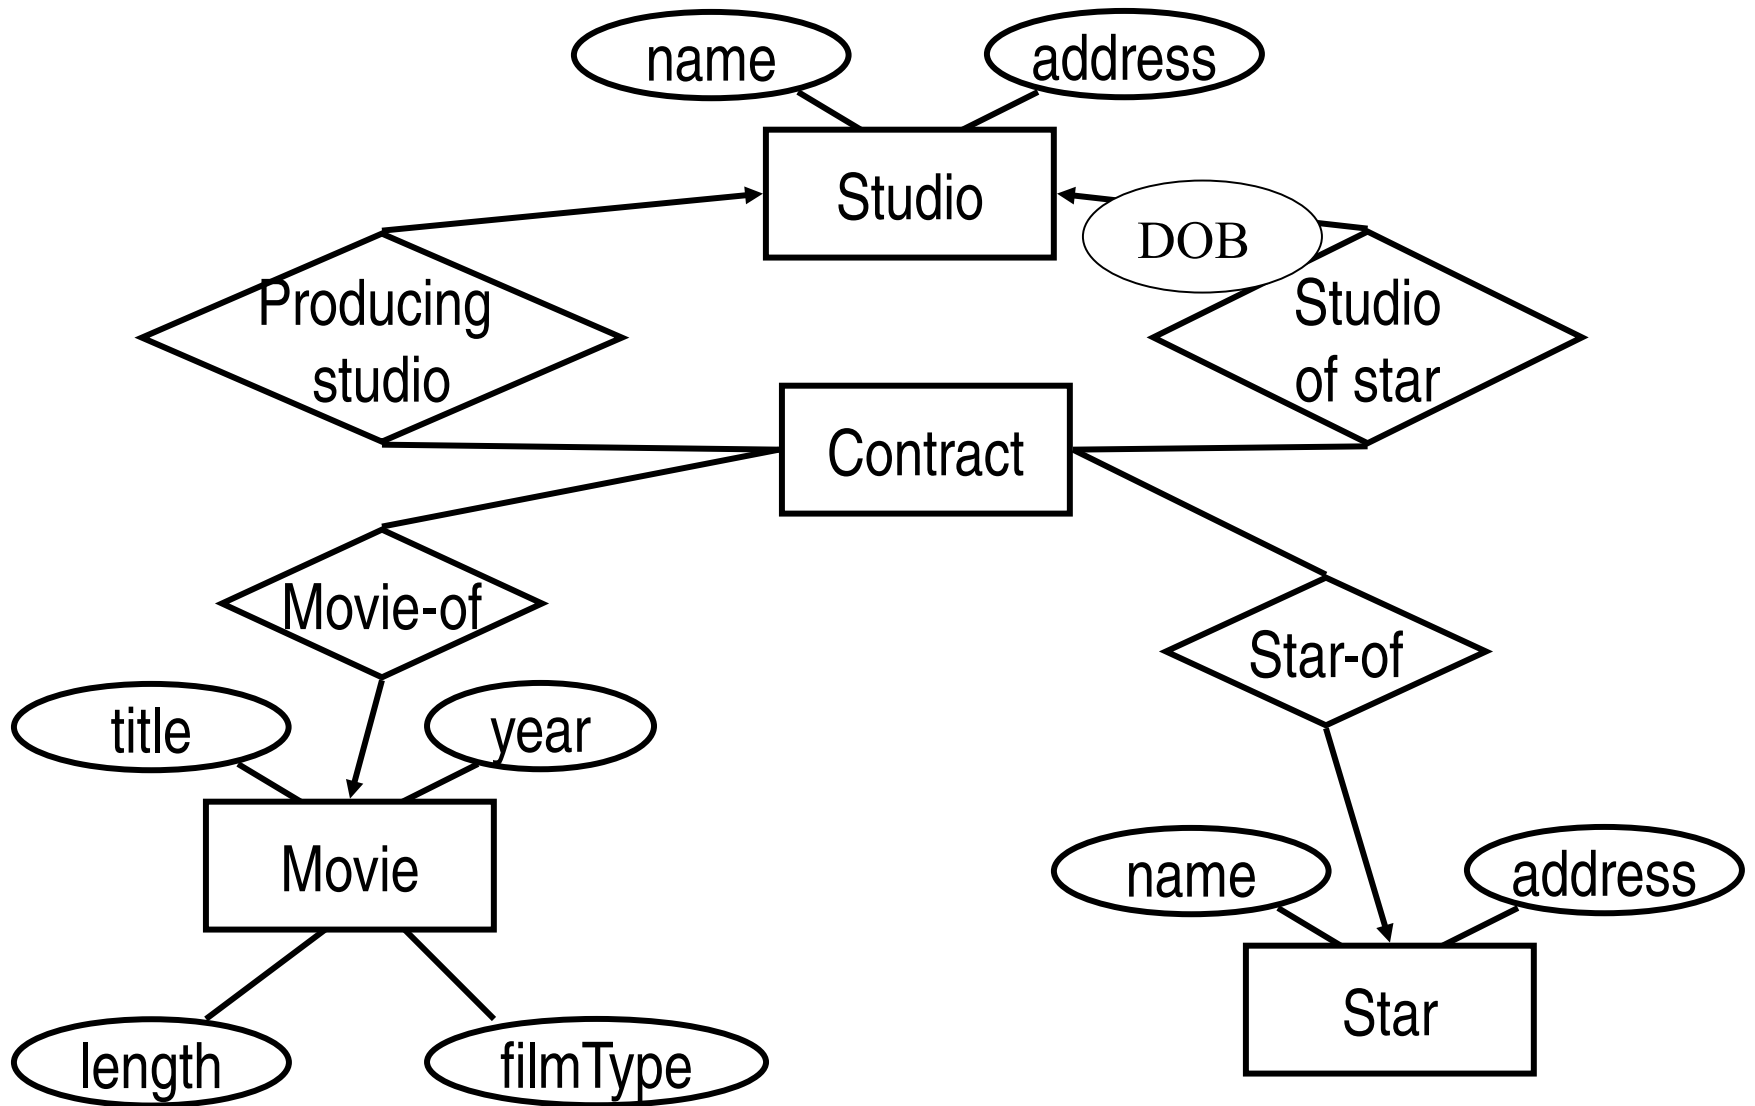
Any n-ary relationship may be converted into a **collection of binary** relationships **without losing** any information???

Introduce a **new** entity set – **connecting** *existing* entity set – whose entities might be thought of as tuples of the relationship for the n-ary relationship

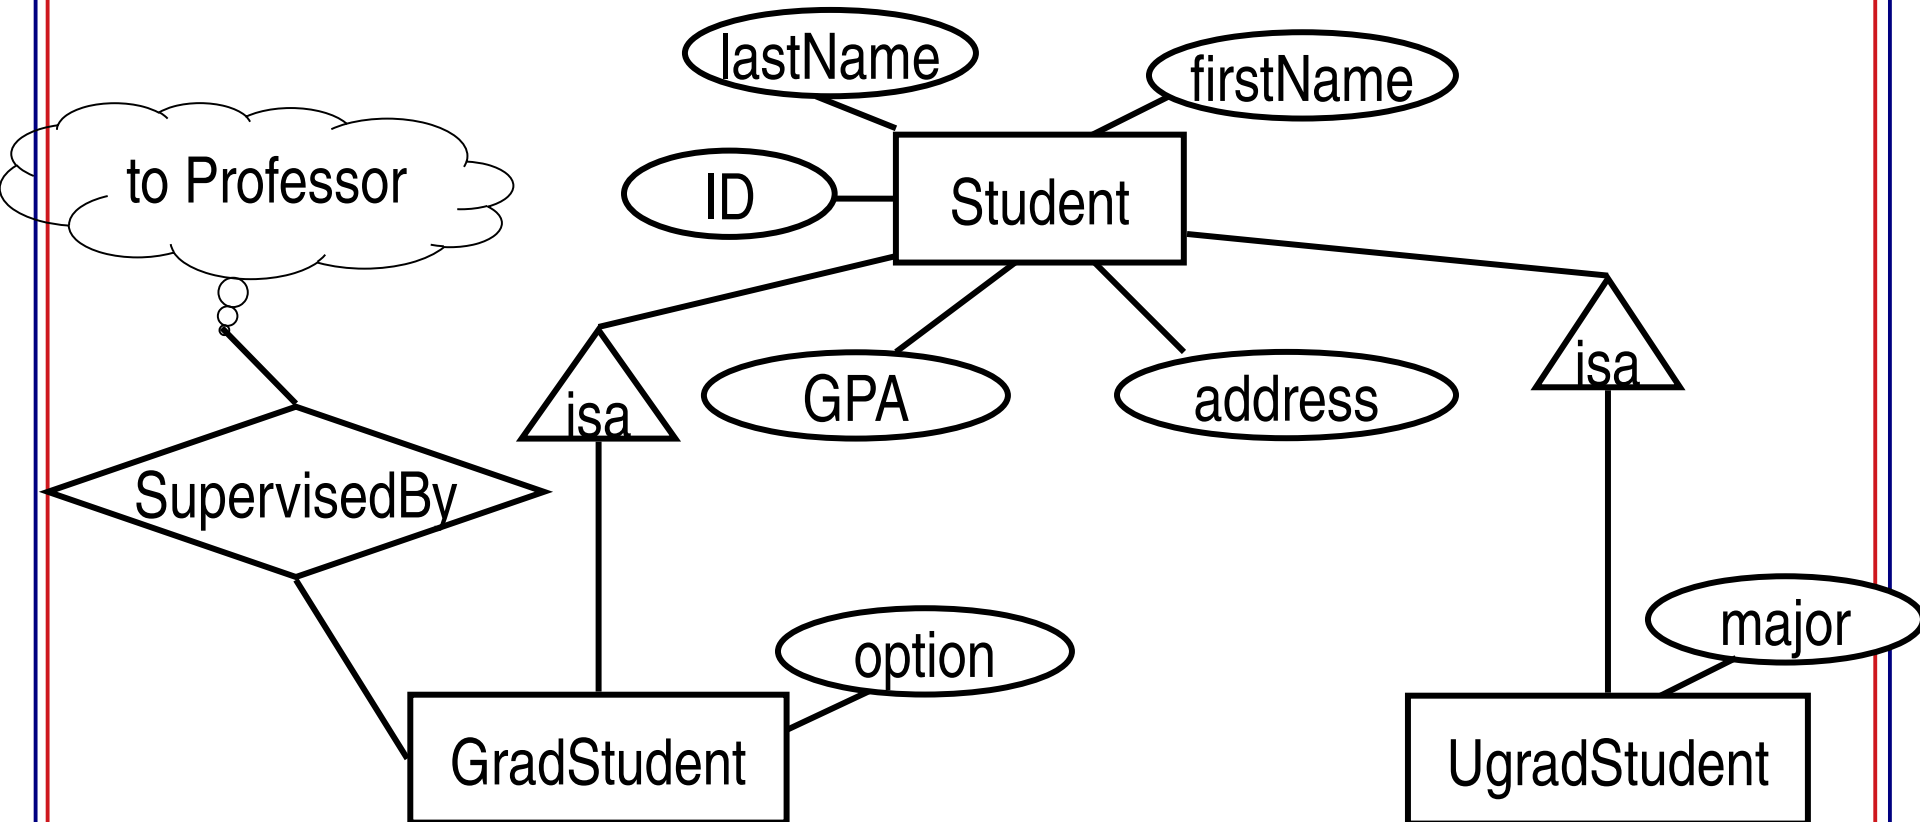
Introduce **many-to-one** relationships from the **connecting** entity set to each of the entity sets participating in the original n-ary relationship

If an entity set plays **more than one** role, then it is the target of one relationship for **each role**

Usually doesn't convey the same semantics – limitation of modeling



Inheritance in E/R is expressed by *isa* relationship



- ❖ There is a subtle difference between the concept of inheritance in ODL and in the E/R model
- ❖ In **ODL**, an object must be a member of exactly one class
- ❖ In the **E/R** model
  - ◆ We shall view an entity as having “components” belonging to several entity sets that are “part of” a single **isa**-hierarchy
  - ◆ The “components” are connected into a single entity by the **isa** relationships
  - ◆ The entity has whatever **attributes** any of its components has, and it participates in whatever **relationships** its components participate in
    - h We need to represent an entity (e.g., CartoonMurderMystery) in the diagram only if it has attributes and/or relationships of its own

# Constraints

- ❖ There are some important **aspects** of the **real world** that **cannot** be represented using the **ODL** or **E/R** model introduced so far
- ❖ The additional information about these aspects often takes the form of **constraints** on the data
- ❖ Sometimes modeling this additional information goes beyond the **structural** and **type** constraints imposed by **classes**, **entity sets**, **attributes**, and **relationships**



**Keys** are (sets of) attributes that **uniquely** identify an object within its class or an entity within its entity set;

**$K \subseteq R$ . *no two entities may agree in all their key values***

**Single-value constraints** are requirements that the value of an attribute be unique. In addition to key constraints, other attributes must have a single-value constraints.

Also in an “one” relationship

**Referential integrity constraints** are requirements that a value referred to by some object must actually exist in the database;

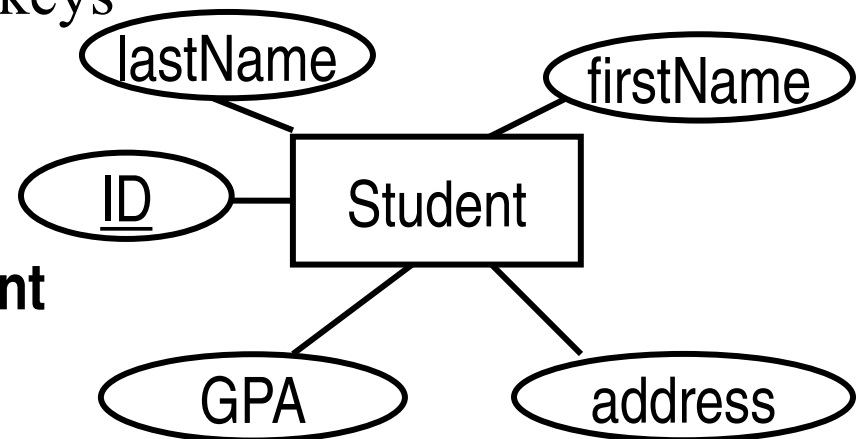
*This means, no dangling pointers.*

**Domain constraints** require that the value of an attribute must be drawn from a specific set of values (called attribute domain), or lies within a specific range

**General constraints** – arbitrary assertions that must hold on the DB

# Keys


- ❖ A *super key* is a set of attributes whose values uniquely identify an entity in the entity set; this set may not be minimal.
- ❖ A **minimal super key** is called a (*candidate*) **key**.
- ❖ An entity may have more than one **key**. One of them is picked as the *primary key*; others may be called *alternate keys*
- ❖ In **E/R**, we underline the key attribute(s) of an entity set (i.e., those attributes that form the primary key of the set)
- ❖ No notation in E/R for alternate keys



ID is the key for the entity set **Student**

# Example


- ❖ What should we select as a key for **Movie** ?
- ❖ **title**?
  - ◆ there could be different movies with the same name
- ❖ **{title, year}**?
  - ◆ there still could be two movies made in the same year, with the same title, but that's very unlikely



**{title, year}** is a key for **Movie**

[https://en.wikipedia.org/wiki/Lists\\_of\\_film\\_remakes](https://en.wikipedia.org/wiki/Lists_of_film_remakes)

# Selecting A Primary Key



Suppose candidate keys for **E** are :

1. {A, B}
2. {D, E}
3. {F, G, H}

❖ Which of the three should “we” pick as the **primary key**?

Criteria to choose a **primary key** when there are more than one candidate:

Total size

Number of attributes

Convenience

A combination of the above

# Single Value Constraint

## ❖ In E/R:

- ◆ attributes are **atomic**
- ◆ an arrow ( $\rightarrow$ ) can be used to express the multiplicity
- ◆ What about multi-valued attributes or relationships?

With the E/R model introduced so far, we cannot express the following options regarding the value of a single-valued **attribute**:

- Require that the value of that attribute to be present(not null)
- Or the presence of the value be optional (null allowed)

If the choice is not explicit, then we may conclude:

- The value must exist if an attribute is part of the key
- The value is optional, otherwise


# Referential Integrity Constraints

❖ For relationships:


◆ **Single-value + Existence = Referential Integrity Constraint**

❖ We extend the **arrow** notation to indicate a reference is mandatory (to support referential integrity)

- The department that gives a course must always exist in the **Department** entity set



# Referential Integrity Constraints



open arrow to denote ref. intg.

- ❖ The studio owning a movie must always be present in the Studios (extent of the Studio entity set)
- ❖ If a president runs a studio, then that studio exists in the Studios
- ❖ There could be studios without a president (temporarily)

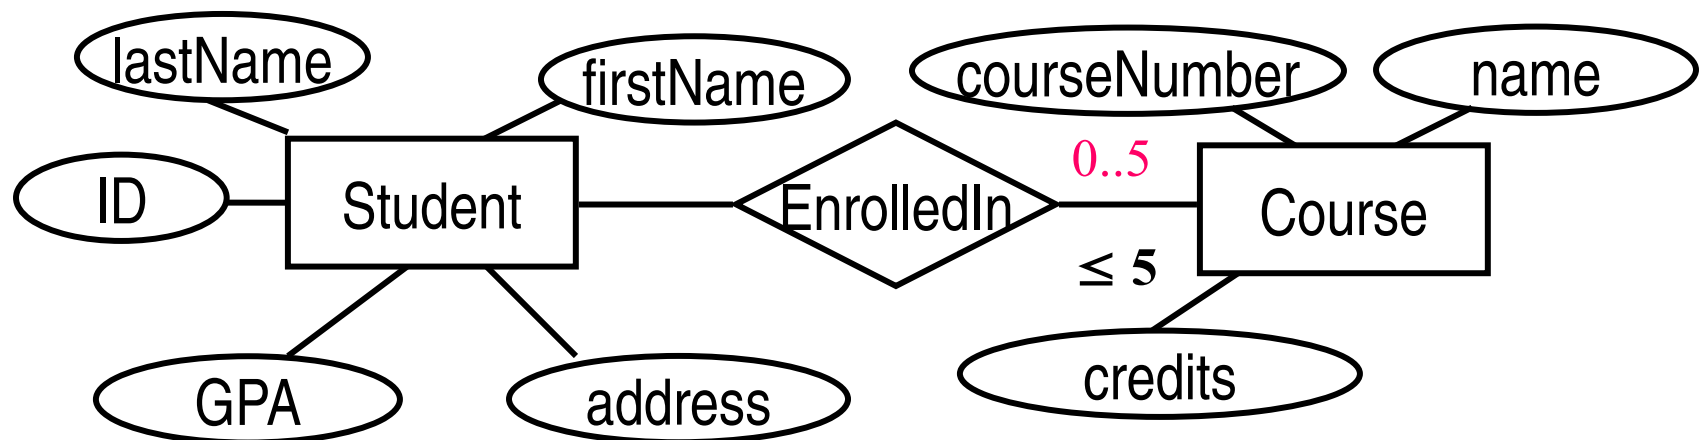
# Domain constraints

- ❖ **Domain constraints** restrict the values of an attribute to be drawn from a set
  - ◆ In ODL, we give a type to the attributes and hence limit their set of values
  - ◆ ODL does **not** support other restrictions, such as that the value should be within a certain range
  - ◆ E/R, in general, does **not** support imposing domain constraints




# Relationship degree constraints

- ❖ **Relationship degree constraints** restrict the number that an entity/object can participate in a relationship
  - ◆ For example, we can impose a constraint saying that a student cannot be enrolled in more than 5 courses
  - ◆ In ODL, we could use, instead of a set of references, an **array** of size 5
  - ◆ In E/R, we may attach a bounding number to the corresponding link




# Weak Entity / Relationship Sets

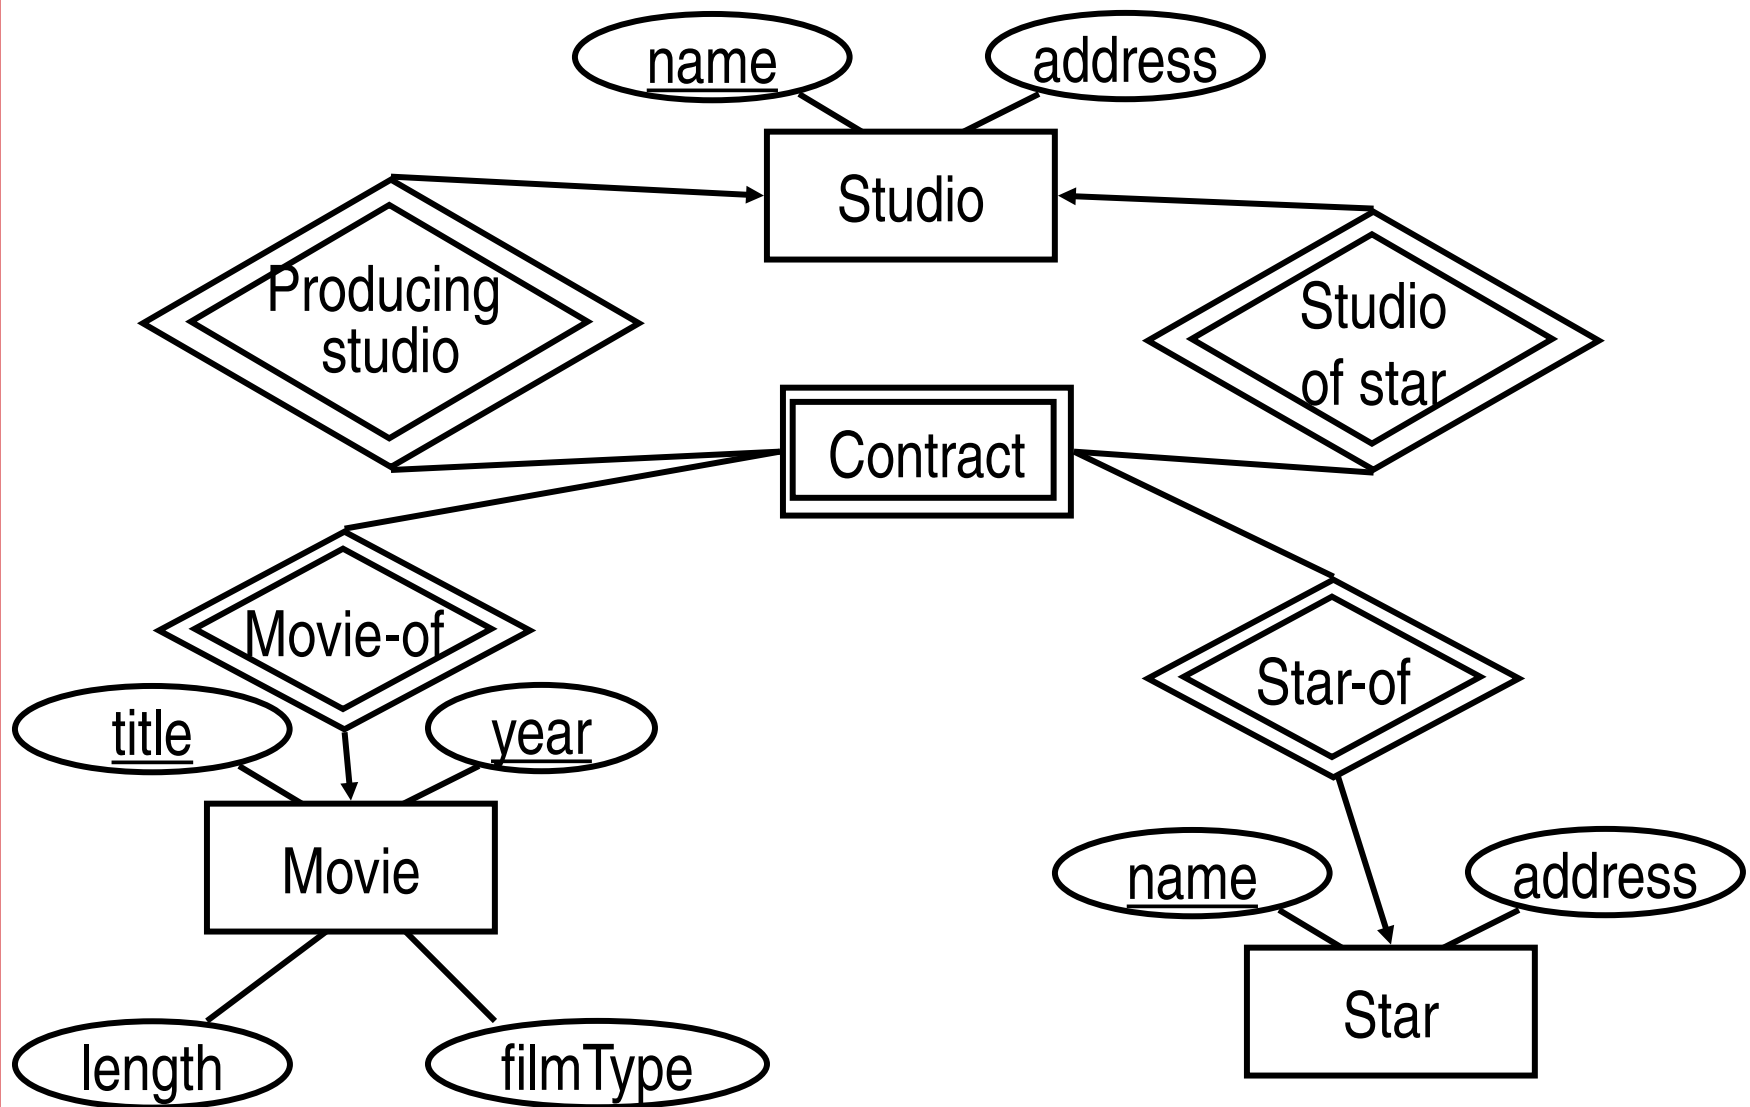
- ❖ A **strong** entity set has a primary key
- ❖ A **weak** entity set does not have sufficient attributes to form a primary key. It should be part of a one-many relationship (with no descriptive attributes) with a strong entity set
- ❖ **Discriminator** of a weak entity set is a set of attributes that distinguishes among the entities corresponding to a strong entity
- ❖ **Primary key** of a weak entity set = primary key of the strong entity + discriminator of the weak entity
- ❖ Represented in E/R model by



# Example

- ❖ Log records transactions done by an ATM
- ❖ Each transaction has a number, date, and an amount
- ❖ Different accounts might have transactions by the same number, on the same date, and for the same amount





# Design Principles

## ❖ Design should

- ◆ Reflect reality

- ◆ Avoid redundancy

  - Redundant information takes space

  - Could cause inconsistency

- ◆ Be as simple as possible

## ❖ Be careful when choosing between using attributes and using classes or entity sets. Remember that

- ◆ An attribute is **simpler** to implement than either a class/entity set or a relationship

- ◆ If something has **more information** associated with it **than just its name**, it probably **needs to be an entity set or a class**

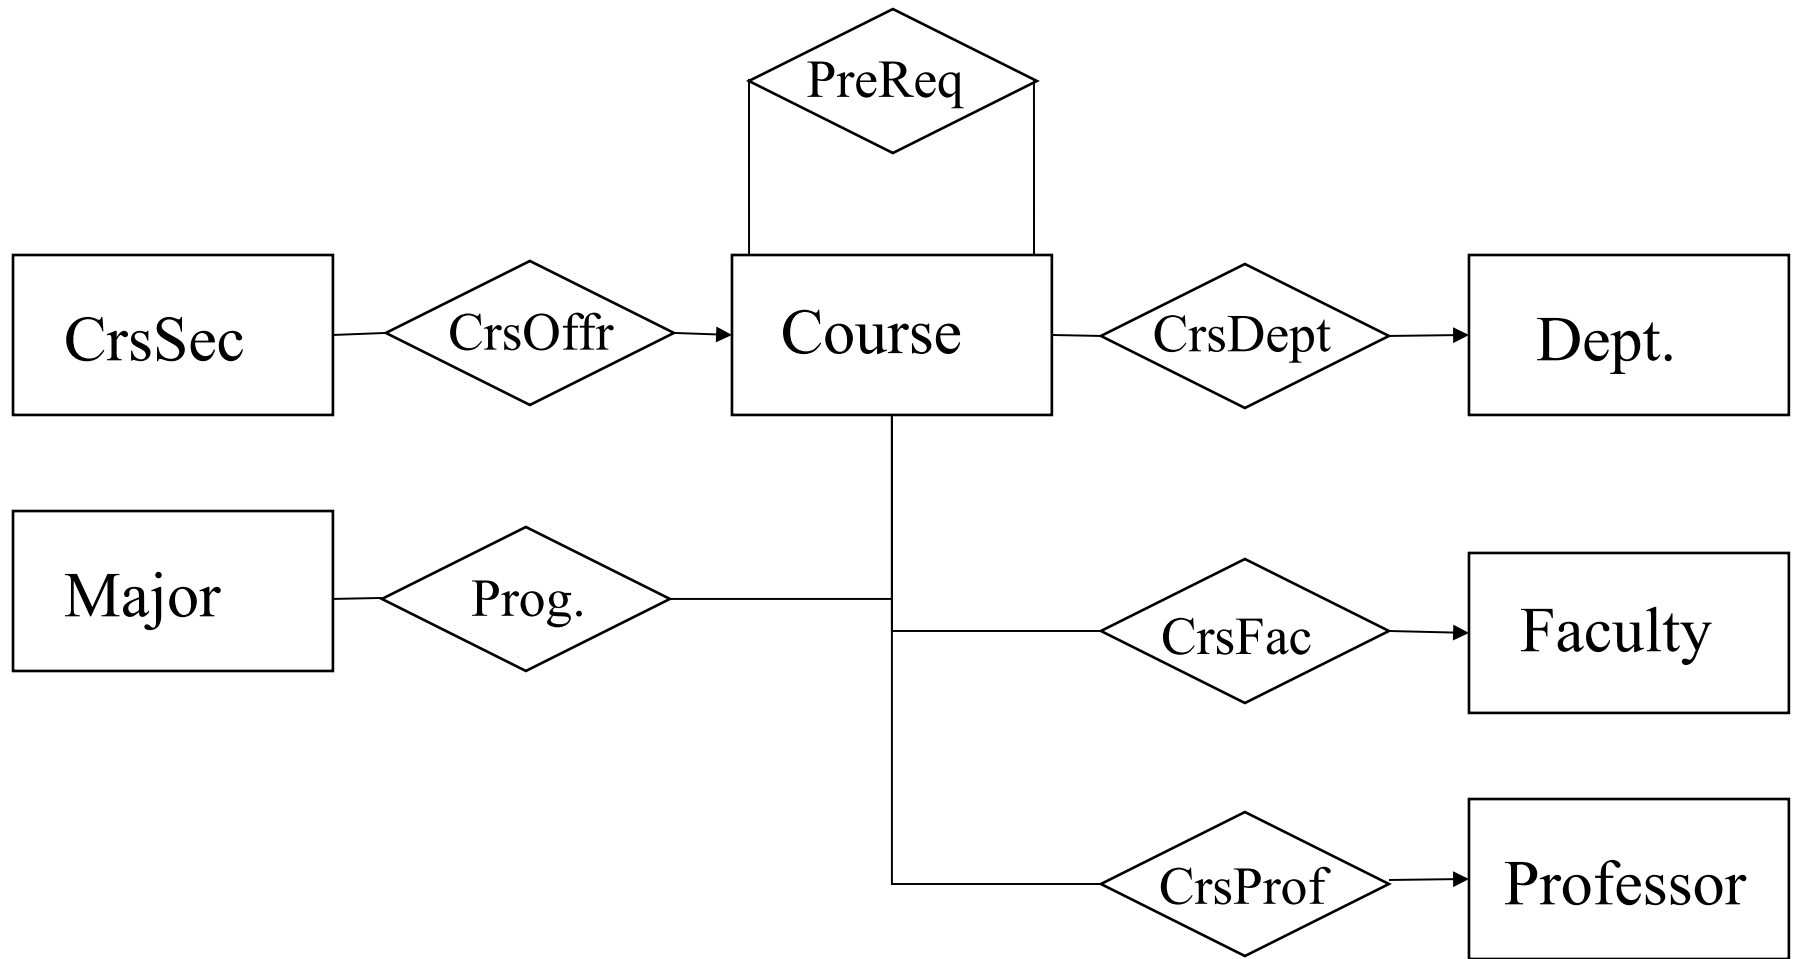
Consider the entity set course in a typical university :

It could be involved in many relationships:

one to many relationships with  
the offering department,  
the offerings faculty  
the professor coordinating the course

A many to one relationship with  
sections for the course

many to many relationships with  
the major program in which it is required  
the pre-requisites (follow up) courses



**How to implement the entity **Course** and its relationship**

# Design decision

Merge the one-to-many relationships  
CrsDept, CrsFac, CrsProf  
in the schema for Course; all attributes of  
these relationships are also included in the  
schema for Course

Create a separate  
relation for each  
one-to-many  
relationships

In this case the relation for Course would have a higher arity;  
but requires one less join to get details for the department, faculty  
or professor for a given course

Similarly, merge the one-to-many relationship  
CrsSec  
in the schema for CrsSec

Create a relation for the many to many relationships  
Program and PreReq